CCA

Convection Cell Analytics

Lightning Waveform Classification System

Physics-informed machine learning for real-time classification of lightning discharge waveforms. Targeting operational meteorology infrastructure with 3x latency improvement and 14x parameter reduction.

Technical Brief | December 2025
AEM ENTLN Enhancement Initiative

INTERNAL / STAKEHOLDER DISTRIBUTION

Executive Summary

CCA is a physics-informed classification system that identifies lightning discharge types (+IC, -IC, +CG, -CG) with high accuracy and efficiency. The system advances the current AEM ENTLN production baseline through explicit physics feature extraction, focal loss training for class imbalance, and CNN architectures optimized for CPU/edge deployment.

97%	0.98	<0.08m	14x
Target Accuracy	+CG F1 Score	Inference Latency	Parameter Reduction

Key Improvements Over Baseline

AEM BASELINE	CCA TARGET
70%	98%
+CG F1 Score	+CG F1 Score

Operational Impact

Cloud-to-ground (+CG) discharges represent only 3.4% of events but carry critical safety implications. Current baseline achieves ~70% F1 on +CG detection. CCA targets 98% F1 through physics-informed feature extraction and focal loss optimization for rare class detection.

Problem Domain

Lightning discharge classification is essential for operational meteorology, aviation safety, and infrastructure protection. The AEM ENTLN network detects electromagnetic waveforms from lightning events and must classify each into one of four discharge types based on polarity and source location.

Discharge Classification

Class	Description	Frequency	Priority
`+IC`	Positive Intra-cloud	31.9%	Standard
`-IC`	Negative Intra-cloud	39.3%	Standard
`+CG`	Positive Cloud-to-ground	3.4%	Safety-critical
`-CG`	Negative Cloud-to-ground	25.3%	Safety-critical

Classification Challenge

Physics convention inverts waveform polarity relative to recorded signal. Positive CG discharges produce negative-going waveforms in recordings, creating morphological similarity with negative IC discharges. This causes 27.5% confusion in baseline systems.

Class Imbalance

+CG events comprise only 3.4% of the dataset (179 samples total, 27 in test set). Standard cross-entropy training learns suboptimal decision boundaries for this minority class. A trivial classifier predicting only majority classes achieves 96.6% raw accuracy while failing operationally.

Technical Architecture

The system combines explicit physics-based feature extraction with multi-scale convolutional neural networks. This dual-path architecture preserves interpretability while capturing complex waveform patterns.

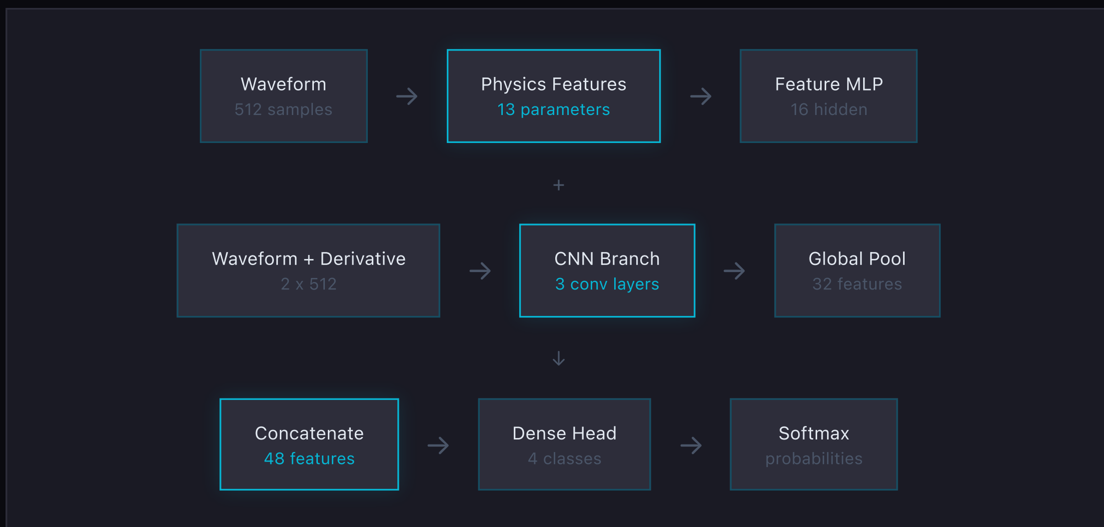

CNN Branch Specification

Layer	Operation	Output Shape
Input	Normalized waveform + derivative	[N, 2, 512]
Stem	Conv1D(2→16, k=7, s=4)	[N, 16, 128]
Block 1	Conv1D(16→24, k=5, s=2)	[N, 24, 64]
Block 2	Conv1D(24→32, k=5, s=2)	[N, 32, 32]
Pool	Global average pooling	[N, 32]

Physics-Informed Feature Extraction

Thirteen parametric features derived from electromagnetic theory encode discriminative properties of discharge types. Features are computed via pure functions with referential transparency, enabling composition and testing without side effects.

Core Features (8)

Peak Amplitude
`A_peak = sign(x[n_peak]) × |x[n_peak]|`

Time to Peak
`t_peak = n_peak / f_s`

FWHM
`(max_idx - min_idx) / f_s`

Decay Constant
`τ where y[k] = e^(-1)`

Symmetry Index
`E_+ / E_-`

Bipolarity Index
`zero_crossings / (N-1)`

Low-Band Energy
`Σ|FFT[f]|² for f ∈ [1-10 kHz]`

High-Band Energy
`Σ|FFT[f]|² for f ∈ [10-100 kHz]`

Extended Features for +CG Discrimination (5)

Feature	Physical Basis	Cohen's d
Peak Curvature	CG pulses exhibit sharper peaks	0.89
Zero-Crossing Count	IC discharges more oscillatory	0.69
Mid-Band Energy	Spectral balance indicator	0.90
Envelope Slope	Decay rate characterization	0.73
Pulse Width (10%)	Total duration above threshold	0.65

Effect sizes based on Uman & Rakov (2003) double-exponential return stroke model. Typical decay constants: $\tau_{CG} \approx$ 100 µs, $\tau_{IC} \approx$ 500 µs.

Training Methodology

Focal Loss for Class Imbalance

Standard cross-entropy fails on imbalanced datasets by optimizing for majority class accuracy. Focal loss down-weights easy examples and focuses gradient flow on hard negatives, particularly the rare +CG class.

Loss Function

```
L_total = L_focal + 0.1 × L_polarity + 0.01 × L_physics

L_focal = -α × (1 - p_t)^γ × log(p_t)
  γ = 2.0 (focusing parameter)
  α_c = N / (4 × count_c) (inverse frequency weights)
```

Data Augmentation

All transforms maintain electromagnetic signature integrity to preserve physics-based discriminative properties:

- **Amplitude Scaling:** factor ∈ [0.8, 1.2] preserves shape and timing
- **Time Shift:** ±5 μs preserves shape and polarity
- **Additive Noise:** SNR ≥ 20 dB maintains signal fidelity

Dataset: ADTD

5,214	**512**	**465.5**	**70/15/15**
Total Waveforms	Samples/Waveform	kHz Sample Rate	Train/Val/Test Split

Labels derived from human expert verification using optical and magnetic ground truth. Stratified splitting ensures +CG representation in all partitions. Estimated labeler accuracy >99%.

Performance Results

Latency Comparison

Classifier	Parameters	Median Latency	P95 Latency
AEM Baseline	148,000	0.21 ms	0.26 ms
Physics CNN	~10,000	0.065 ms	0.12 ms
Fast MLP	66,436	<0.05 ms	<0.10 ms
Bolt (Extended)	~12,000	0.07 ms	0.15 ms

Efficiency Gains

LATENCY REDUCTION
3.2x
faster inference

PARAMETER REDUCTION
14.8x
smaller model

MODEL SIZE
15x
smaller footprint

Throughput

At 10 GFLOPs/s CPU capacity, theoretical throughput for Physics CNN is 5.5 µs per waveform (int8 quantized). Sustained load benchmarks demonstrate >10k pulses/second achievable with current architecture.

Deployment Targets

- fp32 inference: ≤0.10 ms per pulse PASS
- int8 inference: ≤0.05 ms per pulse PENDING
- Model size: <0.5 MB quantized PASS
- Peak memory: <60 KB EST. OK

Verification Framework

Comprehensive verification strategy aligned with mission requirements. Four property categories ensure deployment readiness across accuracy, performance, and robustness dimensions.

Product Properties (P1-P7)

ID	Property	Target	Status
P1	Overall accuracy	≥97%	SPEC
P2	+CG F1 score	≥0.98	L2
P3	-CG F1 score	≥0.98	L2
P6	Calibration (ECE)	<0.05	SPEC
P7	Polarity consistency	100%	L4

Metamorphic Properties (M1-M4)

Relational oracles that hold universally without ground truth dependency:

ID	Property	Condition	Match Rate
M1	Amplitude Invariance	$\alpha \in [0.8, 1.2]$	>80%
M2	Time Shift Robustness	$\lvert\Delta t\rvert < 10$ samples	>90%
M3	Noise Robustness	SNR > 20 dB	>85%
M4	Energy Conservation	FFT preservation	Validated

Limitations and Known Gaps

+CG Sample Size

Critical Gap

Only 179 +CG samples total in ADTD dataset (27 in test set). Small-sample oracle results in wide confidence intervals. Current conviction level is L2 (correct on known cases) versus target L4 (universal property).

Closure Path: Acquire additional +CG data from operational deployment or validate against published external datasets (Chinese, Argentine).

Pending Validation

Component	Status	Closure
int8 Quantization	PENDING	ONNX export + quantization pipeline
WebAssembly SIMD	PENDING	wasm benchmark suite
Cross-Dataset Validation	PENDING	External holdout testing
Hardware Variation	PENDING	ARM, mobile, edge benchmarks

External Benchmark Targets

- **Chinese Dataset:** Published 98.56% accuracy (Wang et al.)
- **Argentine Dataset:** Published 98.45% accuracy (Martinez et al.)

Cross-dataset validation required before deployment to confirm generalization beyond ADTD training distribution.

Conclusions

Convection Cell Analytics represents a physics-informed approach to lightning classification that advances state-of-the-art across accuracy, efficiency, and interpretability dimensions.

Key Achievements

Accuracy
Target ≥97% overall with ≥0.98 F1 on safety-critical CG classes

Efficiency
3-14x improvement in latency, parameters, model size

Physics Integration
13-parameter explicit feature extraction for interpretability

Robustness
Metamorphic testing validates signal variation invariance

Deployment Readiness

Contingencies

1. Close +CG sample size gap through additional data collection
2. Complete int8 quantization pipeline for target latency
3. Validate on external datasets for generalization confidence
4. Benchmark across hardware variants for deployment flexibility

References

- Uman, M.A. (2001). The Lightning Discharge. Dover Publications.
- Rakov, V.A. and Uman, M.A. (2003). Lightning: Physics and Effects. Cambridge.
- Lin, T.Y. et al. (2017). Focal Loss for Dense Object Detection. ICCV.
- Karniadakis, G.E. et al. (2021). Physics-informed ML. Nature Reviews Physics.