

---

# Convection Cell Analytics

Physics-informed lightning waveform classification for edge deployment

TECHNICAL REPORT V0.1.0

PROJECT: CONVECTION-CELL-ANALYTICS

## Overview

---

Convection Cell Analytics (CCA) is a physics-informed machine learning system for real-time classification of lightning discharge waveforms. The system distinguishes four discharge types: positive intra-cloud (+IC), negative intra-cloud (-IC), positive cloud-to-ground (+CG), and negative cloud-to-ground (-CG). Classification accuracy targets 97% overall with  $F1 \geq 0.98$  for safety-critical CG events.

The architecture combines a lightweight 1D CNN with explicit physics feature extraction. Waveforms (512 samples at 465 kHz, spanning 1100  $\mu$ s) are processed through parallel branches: a CNN extracts learned representations while a physics branch computes 13 parametric features encoding electromagnetic domain knowledge. These are fused via concatenation before a classification head.

### DESIGN TENETS

**Physics-Informed Classification:** Leverage electromagnetic signatures inherent in lightning waveforms. **Edge Deployment:** Operate on resource-constrained environments with strict latency (<0.08ms), memory (<60KB), and model size (<0.5MB) constraints.

## Contents

---

|                            |   |
|----------------------------|---|
| 1. Physics Features        | 3 |
| 2. CNN Architecture        | 3 |
| 3. Performance Constraints | 4 |
| 4. Project Structure       | 5 |

## 1. Physics Features

The system extracts 13 parametric features from each waveform, encoding domain knowledge about electromagnetic signatures of different discharge types. Core features (8) provide primary discrimination; extended features (5) improve +CG vs -IC separation.

| FEATURE          | DESCRIPTION                | DISCRIMINATIVE VALUE                                |
|------------------|----------------------------|---|
| Peak Amplitude   | Signed maximum intensity   | Polarity discrimination                             |
| Time to Peak     | Rise time in $\mu\text{s}$ | CG $\sim 5\mu\text{s}$ , IC $\sim 20\mu\text{s}$    |
| FWHM             | Full width at half maximum | Narrow CG, broad IC                                 |
| Decay Constant   | Exponential tail $\tau$    | CG $\sim 100\mu\text{s}$ , IC $\sim 500\mu\text{s}$ |
| Symmetry Index   | E+ / E- energy ratio       | CG unipolar, IC bipolar                             |
| Bipolarity Index | Zero-crossing frequency    | IC oscillatory character                            |
| Energy Bands     | 1-10kHz, 10-100kHz         | Spectral signatures                                 |

Extended features: peak curvature, zero crossing count, mid-band energy (10-50kHz), envelope slope, pulse width.

## 2. CNN Architecture

The network is optimized for CPU inference, avoiding slow dilated convolutions. Three strided convolutions progressively downsample while extracting hierarchical features.

### Network Layers

- Stem: 2 $\rightarrow$ 16 ch, k=7, s=4
- Block 1: 16 $\rightarrow$ 24 ch, k=5, s=2
- Block 2: 24 $\rightarrow$ 32 ch, k=5, s=2
- Global avg pool + Physics MLP
- Head: 48 $\rightarrow$ 4 (softmax)

### Design Choices

- No dilated convolutions
- BatchNorm fused at export
- Focal loss for +CG imbalance
- Total: 6 ops vs 40+ baseline

#### FUSION STRATEGY

```
output = head(concat(cnn, physics))
```

— Physics branch provides interpretable baseline; CNN captures subtle patterns.

### 3. Performance Constraints

Edge deployment requires strict adherence to resource constraints. The following invariants are defined in `mission.yaml` and enforced through automated testing.

#### Classification Accuracy

|                                |                              |
|--------------------------------|------------------------------|
| <b>97%</b><br>OVERALL ACCURACY | <b>≥0.98</b><br>+CG F1 SCORE |
| <b>≥0.98</b><br>-CG F1 SCORE   | <b>4</b><br>CLASSES          |

Cloud-to-ground events are prioritized for safety. F1  $\geq 0.98$  required for both +CG and -CG.

#### Resource Constraints

|                                       |  |
|---------------------------------------|--|
| <b>&lt;0.5MB</b><br>MODEL SIZE (INT8) | <b>&lt;0.08ms</b><br>INFERENCE LATENCY |
| <b>&lt;60KB</b><br>PEAK MEMORY        | <b>&gt;10k/s</b><br>THROUGHPUT         |

Latency ( $<0.08\text{ms}$ ) enables real-time processing. Memory ( $<60\text{KB}$ ) fits microcontrollers. Throughput ( $>10\text{k/s}$ ) handles storm-scale events.

#### Test Matrix

| TEST  | INVARIANT                             | PARAMETERS       |
|-------|---------------------------------------|------------------|
| T1    | Four-class discrimination $\geq 97\%$ | n=400, train=300 |
| T2    | CG F1 threshold $\geq 0.98$           | stratified=true  |
| T3    | Physics features (13)                 | —                |
| T4-T7 | Resource constraints                  | n_iter=100       |
| T8-T9 | Architecture validation               | —                |

## 4. Project Structure

---

The codebase is a UV workspace with two libraries: `waveform` (domain models, physics, classifiers) and `ml-utils` (training infrastructure).

### waveform

- `model.py` — Domain types
- `physics/` — Feature extraction
- `adapters/` — Classifiers
- `data/` — Dataset loaders
- `infra/` — ONNX export

### ml-utils

- `model.py` — Training loop
- `callback.py` — Callbacks
- `tracker/` — Experiment tracking
- `loss.py` — Focal loss
- `benchmark.py` — Performance

## Data Pipeline

Waveforms stored as NumPy arrays (512 samples, float64). Pipeline: normalize (z-score) → derivatives (central diff) → physics features → batch. Masked aggregation handles 1-5 waveforms per pulse.

### DEPENDENCIES

```
numpy, scipy, torch, einops, attrs, expression. Dev: pytest, hypothesis.
```

## Summary

---

CCA combines physics-informed feature extraction with a lightweight CNN for real-time lightning classification on edge devices. The hybrid architecture ensures interpretability through explicit physics features while capturing subtle patterns via learned representations. Strict resource constraints enable deployment on embedded systems and WebAssembly. All invariants are codified in `mission.yaml` and enforced through automated testing.